

EC9170

**DEEP LEARNING FOR ELECTRICAL AND
COMPUTER ENGINEERS**

MINI PROJECT

A.B.O.P. NANDARATHNA - 2021/E/012

A.M. SHANIKA - 2021/E/014

Project Title: "Building an Image Caption Generator using CNN-LSTM from Scratch"

1. Objective

Develop an Image Caption Generator that takes an image as input and generates a meaningful caption using a CNN-LSTM model. The model will be trained from scratch, optimized using hyperparameter tuning and cross-validation, and evaluated to find the bestperforming model.

2. Data Preprocessing Methodology

Dataset Overview

- Source: Kaggle dataset containing 8,090 images and 40,456 captions.
- Structure: Each image has multiple captions.
- Format: Images in .jpg format, captions in a text file.

Image Preprocessing

- Resizing: Images resized to 299x299 (input size for InceptionV3).
- Feature Extraction:
 - Used InceptionV3 (pre-trained on ImageNet) to extract 2048-dimensional feature vectors.
 - Handled grayscale/RGBA images by converting them to RGB.
 - Storage: Extracted features saved in a .pkl file for efficiency.

Text Preprocessing

- Cleaning:
 - Converted to lowercase.
 - Removed punctuation.
 - Stripped extra spaces.
- Tokenization:
 - Used Keras Tokenizer with a vocabulary size of 10,000.
 - Added special tokens: <start>, <end>, <pad>.
- Sequence Preparation:
 - Converted captions to integer sequences.
- Padded sequences to a fixed length of 40 tokens.

3. Model Architecture

Overview

The model follows an encoder-decoder approach:

- Encoder: InceptionV3 (pre-trained CNN) extracts image features.
- Decoder: LSTM-based network generates captions.

Model Components

- Image Feature Processing:
 - Input: 2048-D feature vector from InceptionV3.
 - Passed through a Dense (512 units) + BatchNorm + Dropout layer.
 - Repeated using RepeatVector to match caption length.
- Text Embedding:
 - Input: Padded caption sequences.
 - Embedded into 512-D vectors.
- Sequence Generation:
 - Combined image and text features using Add layer.
 - Processed by LSTM (512 units).
 - Final TimeDistributed Dense layer predicts next word probabilities.
- Key Improvements
 - Teacher Forcing: Trained using shifted sequences.
 - Beam Search: Used during inference for better captions.
 - Regularization: Dropout and BatchNorm to prevent overfitting.

4. Model Training and Optimization

Training Setup

- Optimizer: Adam (lr=0.001).
- Loss: Sparse Categorical Crossentropy (since captions are integers).
- Batch Size: 32.
- Epochs: 20 (with early stopping).

Callbacks

- ModelCheckpoint: Saved best model based on validation loss.
- EarlyStopping: Patience of 5 epochs.
- ReduceLROnPlateau: Reduced learning rate when loss plateaued.

Training Performance

- Training Accuracy: Reached 88.29%.
- Validation Accuracy: 83.87%.
- Loss: Decreased steadily (training: 0.4356, validation: 0.8045).

5. Hyperparameter Tuning and Cross-Validation

Hyperparameters

| Parameter | Value | Notes |
|---------------|--------|--------------------------------|
| Vocab Size | 10,000 | Limited to frequent words |
| Max Length | 40 | Based on caption length stats |
| Embedding Dim | 512 | Higher dimension for semantics |
| LSTM Units | 512 | Balanced capacity & efficiency |
| Dropout | 0.4 | Prevent overfitting |

Cross-Validation

Train-Validation Split: 85%-15%.

Stratified Sampling: Ensured all images had captions in both sets.

6. Model Evaluation and Testing

Evaluation Metrics

| Metric | Score | Interpretation |
|---------|--------|---------------------------------|
| BLEU | 0.0858 | Low but expected for captioning |
| ROUGE-1 | 0.3602 | Moderate word overlap |
| ROUGE-2 | 0.1890 | Captures some bigrams |
| ROUGE-L | 0.3495 | Captures sentence structure |

Qualitative Results

Example 1:

Image: Dog in a park.

Predicted: "White and black dog running in grass."

Actual: "A dog is playing in the park."

Generated Caption: White and black and white dog end.



Example 2:

Image: Beer bottle.

Predicted: "Beer on a table."

Actual: "A bottle of beer with foam."

Generated Caption: Drinks from a beer end.

